

# Serverless Application on AWS

Course Navigation

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

## Monitoring Serverless Application Metrics

Section 6

## Conclusion

Section 7



Linux Academy

# Course Introduction

## Course Introduction

Course Navigation

### Course Introduction

Section 1

#### Course Introduction

About the Training Architect

AWS Free Tier: Usage tracking and Billing Widget

What You Should Know Beforehand:  
Prerequisites

### Serverless Defined

Section 2

### Front-End Serverless Layer: S3 and CloudFront

Section 3

### Application Layer: AWS Lambda Functions

Section 4

### Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5



## Welcome to Fullstack Serverless Application on AWS

Components:

- Front end layer
- Application or Business Logic Layer
- Data persistence Layer

In this course we'll be spending a lot of time on the following:

- AWS CLI API calls
- S3 & S3 Policies
- CloudFront
- Lambda Functions and event triggers
- Aurora RDS Serverless DB
- Linux command line

**Moosa Khalid**  
**AWS Training Architect**

[Back to Main](#)



**Linux Academy**

# Course Introduction

## AWS Free Tier: Usage tracking and Billing Widget

Course Navigation

### Course Introduction

Section 1

Course Introduction

About the Training Architect

AWS Free Tier: Usage tracking and Billing Widget

What You Should Know Beforehand:  
Prerequisites

### Serverless Defined

Section 2

Front-End Serverless Layer: S3 and CloudFront

Section 3

Application Layer: AWS Lambda Functions

Section 4

Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

### 1 Sign up for an AWS Account



User

### 2 Log into the AWS Account



AWS Account

### 3 Set up billing preferences



AWS Management Console



Resources



Alarm



Message

[Back to Main](#)



Linux Academy



Linux Academy

## Course Introduction

Section 1

**Course Introduction**

**About the Training Architect**

**AWS Free Tier: Usage tracking and Billing Widget**

**What You Should Know Beforehand: Prerequisites**

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

[Back to Main](#)

# Course Introduction

## What You Should Know Beforehand: Prerequisites



Before taking this course, here are a few things you should be familiar with:

- AWS services such as **IAM and S3 policies**
- **Python 3 and the Boto AWS SDK** — take our **Automating AWS with Lambda, Python, and Boto3** course to learn more
- **Linux CLI** — we'll execute a lot of Linux CLI commands and AWS CLI commands on Linux OS in the videos and hands-on labs



Python  
(Boto)



Linux Academy

**Course Introduction**

Section 1

**Serverless Defined**

Section 2

**Why Use Serverless?**

Serverless Architectures

Limitations of Serverless

**Front-End Serverless**

Layer: S3 and CloudFront

Section 3

**Application Layer: AWS Lambda Functions**

Section 4

**Back-End Data Persistence Layer: AWS Aurora Serverless**

Section 5

**Monitoring Serverless Application Metrics**

Section 6

[Back to Main](#)

- No servers to **manage**



- Inherent **scaling** and **HA**



- No idle capacity — **lower costs**



- Quick **deployment** and updates



- **Pay-as-you-go** model



- Event/microservices **architectures**

[Next](#)

**Course Introduction**

Section 1

**Serverless Defined**

Section 2

**Why Use Serverless?****Serverless Architectures**

Limitations of Serverless

**Front-End Serverless**Layer: S3 and  
CloudFront

Section 3

**Application Layer: AWS Lambda Functions**

Section 4

**Back-End Data Persistence Layer: AWS Aurora Serverless**

Section 5

**Monitoring Serverless Application Metrics**

Section 6

**Serverless App Components****Front End****Application****Data Persistence**

A **serverless application** doesn't necessarily need all the components defined above to work. It all depends on your **use case and app business logic**.

**Serverless App Architectures****Event-Driven Architecture****Microservices-Based Architecture**

**Serverless services** are excellent for microservices- and event-based architectures, where you can break down the proverbial monolithic monster and focus on **independent, self-contained components** that are excellent at doing one task very efficiently.

In addition to the HA that serverless services inherently offer, having a modular, event-based microservices architecture application on top of it also provides **added stability and disaster recovery**.

**Back****Next****Back to Main****Linux Academy**

**Course Introduction**

Section 1

**Serverless Defined**

Section 2

**Why Use Serverless?****Serverless Architectures**

Limitations of Serverless

**Front-End Serverless****Layer: S3 and CloudFront**

Section 3

**Application Layer: AWS Lambda Functions**

Section 4

**Back-End Data Persistence Layer: AWS Aurora Serverless**

Section 5

**Monitoring Serverless Application Metrics**

Section 6

**Serverless App Components****Front End****Application****Data Persistence**

The **front-end layer** of a serverless ecosystem is basically the presentation layer or the **GUI** (graphical user interface) for the application. A serverless application can just be a **front-end layer** (e.g., an S3 static standalone website, a mobile app's GUI, etc.).

**Serverless App Architectures****Event-Driven Architecture****Microservices-Based Architecture**

**Serverless services** are excellent for microservices- and event-based architectures, where you can break down the proverbial monolithic monster and focus on **independent, self-contained components** that are excellent at doing one task very efficiently.

In addition to the HA that serverless services inherently offer, having a modular, event-based microservices architecture application on top of it also provides **added stability and disaster recovery**.

**Back****Next****Back to Main****Linux Academy**

**Course Introduction**

Section 1

**Serverless Defined**

Section 2

**Why Use Serverless?****Serverless Architectures**

Limitations of Serverless

**Front-End Serverless****Layer: S3 and CloudFront**

Section 3

**Application Layer: AWS Lambda Functions**

Section 4

**Back-End Data Persistence Layer: AWS Aurora Serverless**

Section 5

**Monitoring Serverless Application Metrics**

Section 6

**Serverless App Components****Front End****Application****Data Persistence**

The **application layer** (also known as the business logic layer or middleware) is where the logical decisions of your app are made. In short, **this is where you use "compute" resources**. For example, if you have an online bookshop where users create accounts and buy books, this part of the application will compute what action the user is performing and **direct the inputs and outputs to be stored in a database**, as well as issue receipts and email notifications to users.

**Serverless App Architectures****Event-Driven Architecture****Microservices-Based Architecture**

**Serverless services** are excellent for microservices- and event-based architectures, where you can break down the proverbial monolithic monster and focus on **independent, self-contained components** that are excellent at doing one task very efficiently.

In addition to the HA that serverless services inherently offer, having a modular, event-based microservices architecture application on top of it also provides **added stability and disaster recovery**.

**Back****Next****Back to Main****Linux Academy**

**Course Introduction**

Section 1

**Serverless Defined**

Section 2

**Why Use Serverless?****Serverless Architectures**

Limitations of Serverless

**Front-End Serverless Layer: S3 and CloudFront**

Section 3

**Application Layer: AWS Lambda Functions**

Section 4

**Back-End Data Persistence Layer: AWS Aurora Serverless**

Section 5

**Monitoring Serverless Application Metrics**

Section 6

**Serverless App Components****Front End****Application****Data Persistence**

The application or business logic layer of a serverless app is usually stateless, which means the state of incoming requests or inputs to the function/code is **not saved at the application layer, hence the need for a data persistence layer** (usually some kind of database or data store). For instance, the application code will not remember what product Customer A bought on an online shopping website and hence stores that data persistently to query it at a later time as required. Examples of serverless data layer services include **AWS Aurora Serverless** and **AWS DynamoDB**.

**Event-Driven Architecture****Microservices-Based Architecture**

**Serverless services** are excellent for microservices- and event-based architectures, where you can break down the proverbial monolithic monster and focus on **independent, self-contained components** that are excellent at doing one task very efficiently.

In addition to the HA that serverless services inherently offer, having a modular, event-based microservices architecture application on top of it also provides **added stability and disaster recovery**.

**Back****Next****Back to Main****Linux Academy**

**Course Introduction**

Section 1

**Serverless Defined**

Section 2

**Why Use Serverless?****Serverless Architectures**

Limitations of Serverless

**Front-End Serverless**Layer: S3 and  
CloudFront

Section 3

**Application Layer: AWS Lambda Functions**

Section 4

**Back-End Data Persistence Layer: AWS Aurora Serverless**

Section 5

**Monitoring Serverless Application Metrics**

Section 6

**Serverless App Components****Front End****Application****Data Persistence**

A **serverless application** doesn't necessarily need all the components defined above to work. It all depends on your **use case and app business logic**.

**Serverless App Architectures****Event-Driven Architecture****Microservices-Based Architecture**

- **Actions/code/functions** are triggered in reaction to **events**.
- An event could be a **change in state**.
- Normally, events are triggered via **front-end layer** and consumed by **application and data persistence layers**.
- Allow for **flexible and microservice-based designs**.

**Back****Next****Back to Main****Linux Academy**

**Course Introduction**

Section 1

**Serverless Defined**

Section 2

**Why Use Serverless?****Serverless Architectures**

Limitations of Serverless

**Front-End Serverless**Layer: S3 and  
CloudFront

Section 3

**Application Layer: AWS Lambda Functions**

Section 4

**Back-End Data Persistence Layer: AWS Aurora Serverless**

Section 5

**Monitoring Serverless Application Metrics**

Section 6

**Serverless App Components****Front End****Application****Data Persistence**

A **serverless application** doesn't necessarily need all the components defined above to work. It all depends on your **use case and app business logic**.

**Serverless App Architectures****Event-Driven Architecture****Microservices-Based Architecture**

- Components within application are **loosely coupled**.
- Modular design** — components are independently deployable.
- Components communicate via **lightweight protocols** such as HTTP.
- Works well with **continuous delivery methodologies and event-based architectures**.

**Back****Next****Back to Main****Linux Academy**

**Course Introduction**

Section 1

**Serverless Defined**

Section 2

**Why Use Serverless?****Serverless Architectures****Limitations of Serverless****Front-End Serverless**Layer: S3 and  
CloudFront

Section 3

**Application Layer: AWS Lambda Functions**

Section 4

**Back-End Data Persistence Layer: AWS Aurora Serverless**

Section 5

**Monitoring Serverless Application Metrics**

Section 6

**... And the Limitations**

- Testing and debugging is **challenging**
- Security **concerns**
- Not built for **long-running** processes
- Cold start times can affect resources/functions that are not used frequently, **causing latency**
- Best for applications with short-running processes, as you're charged for the **time your process/code runs**

**Serverless on the Rise**

With an **increasing** amount of users, the serverless community is growing at a fast pace and finding ways to **mitigate or eradicate** all the limits mentioned above as much as possible. For example, with AWS Lambda serverless functions (FaaS), you can now use AWS X-Ray service to debug in depth as well as use **CloudWatch** Events to keep your functions ready and "warmed" up.

[Back](#)[Back to Main](#)

Linux Academy

# Front-End Serverless Layer: S3 and CloudFront

Course Navigation

Overview

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

### Overview

Setting Up AWS S3 as a Standalone Website via Console

AWS CloudFront as an HTTPS Endpoint Provider

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

## Traditional Web Server Front End



- With traditional server "full" architectures, **scaling needs to be monitored** — extra stress
- Servers to **Maintain!**
- Overkill for **single-page**, static website

## Serverless Front End



- S3 scales **automatically** as traffic bursts
- Highly available thanks to being **AWS S3** backed
- No servers to **Maintain** at all
- Developers can easily leverage **SDK/API** for end-to-end setup process

AWS - CloudFront

S3 - Website

Next

Back to Main



Linux Academy

# Front-End Serverless Layer: S3 and CloudFront

Course Navigation

Overview

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

### Overview

Setting Up AWS S3 as a Standalone Website via Console

AWS CloudFront as an HTTPS Endpoint Provider

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

## Traditional Web Server Front End



http/https



WWW



User

- With traditional server "full" architectures, **scaling needs to be monitored** — extra stress
- Servers to **Maintain!**
- Overkill for **single-page**, static website

- Scalable
- Integration with AWS CloudFront
- Easy integration with AWS Route 53 DNS
- For individual web pages or client-side scripts
- Programmatic access via SDK/APIs

AWS - CloudFront

S3 - Website

Next

Back to Main



Linux Academy

# Front-End Serverless Layer: S3 and CloudFront

Course Navigation

Overview

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

### Overview

Setting Up AWS S3 as a Standalone Website via Console

AWS CloudFront as an HTTPS Endpoint Provider

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

- Content delivery network of edge locations
- Caches static files for lower latency
- DDos protection
- Offloads website
- Integrates well with AWS S3
- Offers HTTPS termination via ACM
- Programmatic access via SDK/APIs

## Serverless Front End



- S3 scales **automatically** as traffic bursts
- Highly available thanks to being **AWS S3** backed
- No servers to **Maintain** at all
- Developers can easily leverage **SDK/API** for end-to-end setup process

AWS - CloudFront

S3 - Website

Next

Back to Main



Linux Academy

# Front-End Serverless Layer: S3 and CloudFront

Course Navigation

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

### Overview

### Setting Up AWS S3 as a Standalone Website via Console

AWS CloudFront as an HTTPS Endpoint Provider

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

### Setting Up AWS S3 as a Standalone Website via Console



S3 Bucket  
with Objects



S3 Static  
Website

- **Static website:** A website with fixed content. Shows the same content to all viewers.
- **S3 static website bucket** is simply a bucket with a property enabled to allow it to host content.
- **Remember:** S3 website URL is not the same as an S3 bucket URL!
- Content hosted by a website bucket has to be **publicly readable** — i.e., `s3:GetObject` allowed on it.
- You cannot have **server-side code execution** since there's no back-end servers to run code on. Only **client-side scripting and individual web pages** can be hosted.
- **S3 policies** and **ACLs** still apply here, so beware of allowing read access on your publicly visible content within the bucket.

Configuring S3 via CLI

Back

Next

Back to Main



Linux Academy

# Front-End Serverless Layer: S3 and CloudFront

Course Navigation

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

### Overview

#### Setting Up AWS S3 as a Standalone Website via Console

AWS CloudFront as an HTTPS Endpoint Provider

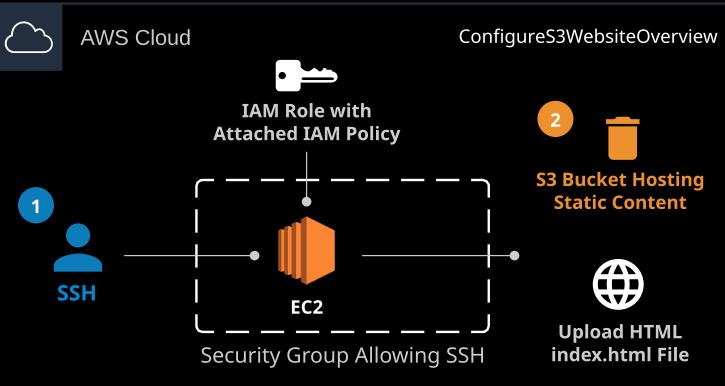
## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

### Setting Up AWS S3 as a Standalone Website via Console



### Summary Breakdown of Steps

1 Have AWS CLI API installed. In our case, we log in via SSH to an EC2 instance already set up with it.

#### Intermediate but important step:

Create and attach an IAM role (allowing S3 read/write access) to your EC2 instance so it can assume that role and make successful API calls.

2 Issue the correct AWS CLI commands to create and modify S3 bucket and upload your website files in the most basic instance, namely: index.html.

#### Supports two URL formats:

1. <bucket-name>.s3-website.<AWS-region>.amazonaws.com
2. <bucket-name>.s3-website-<AWS-region>.amazonaws.com
  - S3 static website requires public read access on bucket
  - Does not support HTTPS

Back

Back to Main



Linux Academy

# Front-End Serverless Layer: S3 and CloudFront

Course Navigation

AWS CloudFront as an HTTPS Endpoint Provider

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

### Overview

Setting Up AWS S3 as a Standalone Website via Console

AWS CloudFront as an HTTPS Endpoint Provider

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

CloudFront via CLI



Amazon  
CloudFront

- Worldwide group of **edge** (caching) locations.
- On initial creation, a **CloudFront distribution** takes about 20 minutes to sync across the globe.
- Works via "**web distributions**" in our scenario, which are hooked to the origin (server, bucket, EC2 instance, etc.) from where your content is to be served and cached from.
- User can change value of "**viewer protocol policy**" in web distribution to redirect all HTTP traffic to HTTPS.
- CloudFront provides default **SSL certificates** you can use for **HTTPS**; however, users can upload their own custom certs and route to CloudFront from a **DNS** such as Route 53.

Low-Latency  
Access to Content  
via Caching

DDoS  
Prevention

HTTPS  
Endpoint

Back

Back to Main



Linux Academy

# Front-End Serverless Layer: S3 and CloudFront

Course Navigation

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

### Overview

Setting Up AWS S3 as a Standalone Website via Console

AWS CloudFront as an HTTPS Endpoint Provider

## Application Layer: AWS Lambda Functions

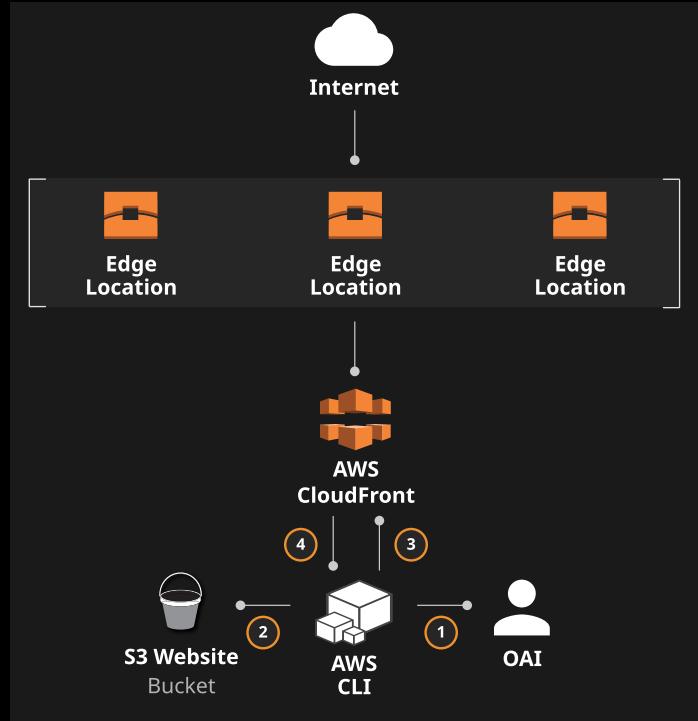
Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

### AWS CloudFront as an HTTPS Endpoint Provider

CloudFront via CLI



Back

Back to Main



Linux Academy

# Front-End Serverless Layer: S3 and CloudFront

Course Navigation

AWS CloudFront as an HTTPS Endpoint Provider

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

### Overview

Setting Up AWS S3 as a Standalone Website via Console

AWS CloudFront as an HTTPS Endpoint Provider

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

CloudFront via CLI



Amazon  
CloudFront

- Worldwide group of **edge** (caching) locations.
- On initial creation, a **CloudFront distribution** takes about 20 minutes to sync across the globe.
- Works via "**web distributions**" in our scenario, which are hooked to the origin (server, bucket, EC2 instance, etc.) from where your content is to be served and cached from.
- User can change value of "**viewer protocol policy**" in web distribution to redirect all HTTP traffic to HTTPS.
- CloudFront provides default **SSL certificates** you can use for **HTTPS**; however, users can upload their own custom certs and route to CloudFront from a **DNS** such as Route 53.

Low-Latency  
Access to Content  
via Caching

DDoS  
Prevention

HTTPS  
Endpoint

1

Create an **origin access identity** in CloudFront. It is a special CloudFront user that can access an **S3 bucket on behalf of users** accessing objects in a website bucket via the CloudFront service.

BACK

Back to Main



Linux Academy

# Front-End Serverless Layer: S3 and CloudFront

Course Navigation

AWS CloudFront as an HTTPS Endpoint Provider

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

### Overview

Setting Up AWS S3 as a Standalone Website via Console

AWS CloudFront as an HTTPS Endpoint Provider

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

CloudFront via CLI



Amazon  
CloudFront

- Worldwide group of **edge** (caching) locations.
- On initial creation, a **CloudFront distribution** takes about 20 minutes to sync across the globe.
- Works via "**web distributions**" in our scenario, which are hooked to the origin (server, bucket, EC2 instance, etc.) from where your content is to be served and cached from.
- User can change value of "**viewer protocol policy**" in web distribution to redirect all HTTP traffic to HTTPS.
- CloudFront provides default **SSL certificates** you can use for **HTTPS**; however, users can upload their own custom certs and route to CloudFront from a **DNS** such as Route 53.

Low-Latency  
Access to Content  
via Caching

DDoS  
Prevention

HTTPS  
Endpoint

2

Attach an **S3 bucket policy** to the already created S3 bucket website to allow CloudFront OAI user **s3:GetObject** access on all objects within the bucket.

Back

Back to Main



Linux Academy

# Front-End Serverless Layer: S3 and CloudFront

Course Navigation

AWS CloudFront as an HTTPS Endpoint Provider

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

### Overview

Setting Up AWS S3 as a Standalone Website via Console

AWS CloudFront as an HTTPS Endpoint Provider

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

CloudFront via CLI



Amazon  
CloudFront

- Worldwide group of **edge** (caching) locations.
- On initial creation, a **CloudFront distribution** takes about 20 minutes to sync across the globe.
- Works via "**web distributions**" in our scenario, which are hooked to the origin (server, bucket, EC2 instance, etc.) from where your content is to be served and cached from.
- User can change value of "**viewer protocol policy**" in web distribution to redirect all HTTP traffic to HTTPS.
- CloudFront provides default **SSL certificates** you can use for **HTTPS**; however, users can upload their own custom certs and route to CloudFront from a **DNS** such as Route 53.

Low-Latency  
Access to Content  
via Caching

DDoS  
Prevention

HTTPS  
Endpoint

3

Create a **CloudFront** distribution using **CLI**. It caches and serves your content via edge locations.

Back

Back to Main



Linux Academy

# Front-End Serverless Layer: S3 and CloudFront

Course Navigation

AWS CloudFront as an HTTPS Endpoint Provider

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

### Overview

Setting Up AWS S3 as a Standalone Website via Console

AWS CloudFront as an HTTPS Endpoint Provider

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

CloudFront via CLI



### Amazon CloudFront

- Worldwide group of **edge** (caching) locations.
- On initial creation, a **CloudFront distribution** takes about 20 minutes to sync across the globe.
- Works via "**web distributions**" in our scenario, which are hooked to the origin (server, bucket, EC2 instance, etc.) from where your content is to be served and cached from.
- User can change value of "**viewer protocol policy**" in web distribution to redirect all HTTP traffic to HTTPS.
- CloudFront provides default **SSL certificates** you can use for **HTTPS**; however, users can upload their own custom certs and route to CloudFront from a **DNS** such as Route 53.

Low-Latency Access to Content via Caching

DDoS Prevention

HTTPS Endpoint

4

Get the configuration of the **CloudFront distribution you just created**, modify it (PriceClass, OAI, redirect all to HTTPS), and update the original CloudFront distribution configuration with your **modified one — all via CLI**.

BACK

Back to Main



Linux Academy

# Application Layer: AWS Lambda Functions

Course Navigation

Overview

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

### Overview

Creating a Lambda Function and Lambda IAM Execution Role

Event Triggers for AWS Lambda

Testing and Debugging Lambda Functions

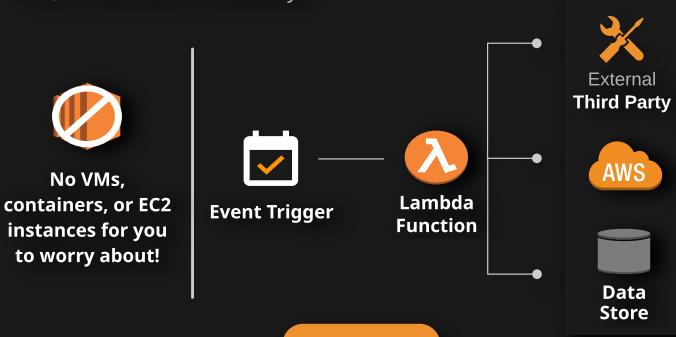
## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

### Under the Hood: Lambda

**Function as a Service** (FaaS) is a type of serverless, event-driven computing.

**No containers or VMs** are present, and users just focus on the code. Plus, it has built-in scalability.



- Charged **based on requests + duration of execution**

- Supports **blue/green** deployments
- Allows useful triggers from **S3 and SNS** to execute
- A single dial configuration of memory, with which **CPU scales proportionally**
- Gels with **CloudFront** to offer Lambda@Edge
- Automatic scaling, **built-in fault tolerance**
- Bring your own code or language using **Lambda Layers**

Back

Next

Back to Main



Linux Academy

# Application Layer: AWS Lambda Functions

Course Navigation

Overview

## Course Introduction

Section 1

Under the Hood: Lambda

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

### Overview

Creating a Lambda Function and Lambda IAM Execution Role

Event Triggers for AWS Lambda

Testing and Debugging Lambda Functions

Back

There really is no such thing as "serverless." Rather, it's someone else's fully managed execution environment you use at a fraction of the cost. Lambda function containers are like the Oompa Loompas in *Charlie and the Chocolate Factory* — always working and available in the back end but never seen or interacted with.



## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

Back

Next

Back to Main



Linux Academy

# Application Layer: AWS Lambda Functions

Course Navigation

Overview

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

### Overview

Creating a Lambda Function and Lambda IAM Execution Role

Event Triggers for AWS Lambda

Testing and Debugging Lambda Functions

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

### Under the Hood: Lambda

**Function as a Service (FaaS)** is a type of serverless, event-driven computing.  
**Event trigger comes in**

**No containers or VMs** are present, and users just focus on the code.  
Plus, it has built-in scalability.



Back

Next

Back to Main



Linux Academy

**Course Introduction**

Section 1

**Serverless Defined**

Section 2

**Front-End Serverless Layer: S3 and CloudFront**

Section 3

**Application Layer: AWS Lambda Functions**

Section 4

**Overview****Creating a Lambda Function and Lambda IAM Execution Role**

Event Triggers for AWS Lambda

Testing and Debugging Lambda Functions

**Back-End Data Persistence Layer: AWS Aurora Serverless**

Section 5

[Back to Main](#)

# Application Layer: AWS Lambda Functions

## Creating a Lambda Function and Lambda IAM Execution Role

IAM  
Role

Code

Runtime  
ConfigurationLambda  
Function

### 1. Recipe for Lambda function creation:

- Provide the **runtime or language** you want to code in.
- Provide an **IAM role Lambda** can assume to carry out tasks, such as creating logs in CloudWatch Logs and streaming to the logs in **CloudWatch** log groups created by it.

### 2. Other options set by default but can be modified:

- Memory size (one dial for controlling it all)
- Concurrency** across AWS account
- Env variables/debugging (SQS DLQ/AWS X-Ray)
- ZIP file or **S3** location of your code
- Add triggers to **invoke the function**

## Lambda IAM Execution Role

The **Lambda service** requires an IAM role with specific permissions so it can query or reach out to other AWS services on your behalf through your code. It's best practice to allow least possible access to any resource within AWS. At a minimum, Lambda requires **access to CloudWatch Logs for log streaming**.

[Console vs. CLI: Lambda](#)[Back](#)[Next](#)

Linux Academy

**Course Introduction**

Section 1

**Serverless Defined**

Section 2

**Front-End Serverless Layer: S3 and CloudFront**

Section 3

**Application Layer: AWS Lambda Functions**

Section 4

**Overview****Creating a Lambda Function and Lambda IAM Execution Role**

Event Triggers for AWS Lambda

Testing and Debugging Lambda Functions

**Back-End Data Persistence Layer: AWS Aurora Serverless**

Section 5

# Application Layer: AWS Lambda Functions

## Creating a Lambda Function and Lambda IAM Execution Role

IAM  
Role

Code

Runtime  
ConfigurationLambda  
Function

### 1. Recipe for Lambda function creation:

- Provide the **runtime or language** you want to code in.
- Provide an **IAM role Lambda** can assume to carry out tasks, such as creating logs in CloudWatch Logs and streaming to the logs in **CloudWatch** log groups created by it.

### 2. Other options set by default but can be modified:

- Memory size (one dial for controlling it all)
- Concurrency** across AWS account
- Env variables/debugging (SQS DLQ/AWS X-Ray)
- ZIP file or **S3** location of your code
- Add triggers to **invoke the function**

A quick note about creating Lambda functions via the **AWS Management Console** as opposed to the **AWS CLI**: The AWS console makes it easy and seamless to create a Lambda role on the fly; whereas with the CLI, you need to **create a Lambda service IAM role** and attach the relevant IAM policy with it to then pass into your **create-function** invocation.

**Console vs. CLI: Lambda****Back****Next****Back to Main****Linux Academy**

# Application Layer: AWS Lambda Functions

Course Navigation

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

### Overview

### Creating a Lambda Function and Lambda IAM Execution Role

### Event Triggers for AWS Lambda

Testing and Debugging Lambda Functions

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

## Event Triggers for AWS Lambda

### Events



**Basically, integrations with other services to invoke Lambda functions.**

Each service that integrates with Lambda **sends data to your Lambda function in JSON** as an event, which is passed as an "event" object into your handler function. It can be set from within **Lambda functions** or the **CloudWatch events console**.

### Lambda can also:

- Read events from services such as: **Kinesis, DynamoDB and SQS**
- Be invoked by services **synchronously**
- Be invoked by services **asynchronously**

### Synchronous Events

### Asynchronous Events

[Back to Main](#)

[Back](#)

[Next](#)



Linux Academy

# Application Layer: AWS Lambda Functions

Course Navigation

## Event Triggers for AWS Lambda

### Course Introduction

Section 1

### Serverless Defined

Section 2

### Front-End Serverless Layer: S3 and CloudFront

Section 3

### Application Layer: AWS Lambda Functions

Section 4

#### Overview

#### Creating a Lambda Function and Lambda IAM Execution Role

#### Event Triggers for AWS Lambda

Testing and Debugging Lambda Functions

### Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

Events



Event triggers can be from a number of sources, including S3, IoT, CloudFront, DynamoDB, and even external third-party tools, such as Datadog and PagerDuty via CloudWatch event bus and AWS EventBridge services.

#### Lambda can also:

- Read events from services such as: **Kinesis, DynamoDB and SQS**
- Be invoked by services **synchronously**
- Be invoked by services **asynchronously**

Synchronous Events

Asynchronous Events

Back

Next

Back to Main



Linux Academy

# Application Layer: AWS Lambda Functions

Course Navigation

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

### Overview

### Creating a Lambda Function and Lambda IAM Execution Role

### Event Triggers for AWS Lambda

Testing and Debugging Lambda Functions

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

## Event Triggers for AWS Lambda

### Events



**Basically, integrations with other services to invoke Lambda functions.**

Each service that integrates with Lambda **sends data to your Lambda function in JSON** as an event, which is passed as an "event" object into your handler function. It can be set from within **Lambda functions** or the **CloudWatch events console**.

### Lambda can also:

- Read events from services such as: **Kinesis, DynamoDB and SQS**
- Be invoked by services **synchronously**
- Be invoked by services **asynchronously**

### Synchronous Events

### Asynchronous Events

Synchronous services invoke Lambda, and then wait for a response or retry on failure. Examples of such services are **AWS ELB, CloudFront, and API Gateway**.

Back

Next

[Back to Main](#)



Linux Academy

# Application Layer: AWS Lambda Functions

Course Navigation

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

### Overview

### Creating a Lambda Function and Lambda IAM Execution Role

### Event Triggers for AWS Lambda

Testing and Debugging Lambda Functions

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

### Event Triggers for AWS Lambda

Events



**Basically, integrations with other services to invoke Lambda functions.**

Each service that integrates with Lambda **sends data to your Lambda function in JSON** as an event, which is passed as an "event" object into your handler function. It can be set from within **Lambda functions** or the **CloudWatch events console**.

**Lambda can also:**

- Read events from services such as: **Kinesis, DynamoDB and SQS**
- Be invoked by services **synchronously**
- Be invoked by services **asynchronously**

Synchronous Events

Asynchronous Events

For **asynchronous** services, Lambda queues the events before passing it to the function, and the service gets a success response as soon as the event is queued. The async service **doesn't care what happens afterwards**. Examples of such services are Amazon S3, SNS, SES, and CloudFormation.

Back

Next

[Back to Main](#)



Linux Academy

# Application Layer: AWS Lambda Functions

Course Navigation

## Course Introduction

Section 1

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

### Overview

### Creating a Lambda Function and Lambda IAM Execution Role

### Event Triggers for AWS Lambda

### Testing and Debugging Lambda Functions

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

## Testing and Debugging Lambda Functions



AWS  
X-Ray



AWS SAM  
CLI



AWS Cloud9



AWS  
CloudTrail



Lambda  
Function



Amazon  
CloudWatch  
Logs

**AWS has been introducing new ways to test and debug Lambda functions:**

- In the cloud using CloudWatch Logs, AWS X-Ray, and Cloud9 IDE
  - Locally via AWS SAM CLI
  - In the Lambda AWS console directly with test JSON data (can also be done via AWS CLI)
- 
- You can use **logging statements** within your code to help troubleshoot.
  - Investigate issues using **AWS CloudWatch metrics** for Lambda.

Back

Back to Main



Linux Academy

# Back-End Data Persistence Layer: AWS Aurora Serverless

Course Navigation

Overview

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

### Overview

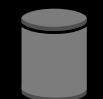
What Is Amazon Aurora Serverless?

Aurora Serverless Concepts

Integrating Aurora Serverless with Lambda Functions

## Monitoring Serverless Application Metrics

Section 6



Database

- **Amazon Lambda** does not offer data persistence, and it does not save states between transactions, which is why a persistent data resource is required.
- **Lambda** can write out to a persistent resource.

Candidates for persistent layer:

- **Amazon S3**
- **Amazon RDS (Aurora Serverless)**
- **Amazon DynamoDB**

Continuance of **data existence** after the process that created it has terminated.

### INSERT/UPDATE/DELETE/CREATE



Lambda Function



AWS Aurora RDBMS

The Lambda function code connects with and carries out **modification actions** on the Aurora Serverless **RDBMS**.

Next

Back to Main



Linux Academy

# Back-End Data Persistence Layer: AWS Aurora Serverless

Course Navigation

What Is Amazon Aurora Serverless?

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

### Overview

#### What Is Amazon Aurora Serverless?

Aurora Serverless Concepts

Integrating Aurora Serverless with Lambda Functions

## Monitoring Serverless Application Metrics

Section 6

### SERVLESS SQL RELATIONAL DATABASE

Managed AutoScaling

Fully AWS Managed

Pay as You Go

### Features

Aurora Serverless manages a **warm pool of resources** in an AWS region to **minimize scaling time**.

You're billed in **Aurora capacity units (ACUs)**, which is a combination of processing and memory capacity. In creating Aurora Serverless, you set a minimum and maximum **ACU** that it will be scaled between.

Like other RDS databases, users receive a **database endpoint** after provisioning. The difference is there's no need to specify **DB instance size or class** when creating Aurora Serverless instances.

Back

Next

Back to Main



Linux Academy

# Back-End Data Persistence Layer: AWS Aurora Serverless

Course Navigation

What Is Amazon Aurora Serverless?

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

### Overview

### What Is Amazon Aurora Serverless?

Aurora Serverless Concepts

Integrating Aurora Serverless with Lambda Functions

## Monitoring Serverless Application Metrics

Section 6

### SERVERLESS SQL RELATIONAL DATABASE

Managed AutoScaling

Fully AWS Managed

Pay as You Go

### Features

- MySQL and PostgreSQL compatible.
- DB storage automatically varies between 10 GiB to 64 TiB.
- DB cluster parameter groups can be used to modify DB instance parameters.
- If there are no connections or transactions coming in, Aurora Serverless billing is paused. You only pay for the storage.
- By default, service is paused after five minutes of inactivity (can be modified).

Like other RDS databases, users receive a **database endpoint** after provisioning. The difference is there's no need to specify **DB instance size or class** when creating Aurora Serverless instances.

Back

Next

[Back to Main](#)



Linux Academy

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

### Overview

### What Is Amazon Aurora Serverless?

### Aurora Serverless Concepts

Integrating Aurora Serverless with Lambda Functions

## Monitoring Serverless Application Metrics

Section 6

Aurora Serverless: Behind the Scenes



Aurora Serverless

## Why RDS Aurora Serverless?

- Provisioning for **peak** capacity ==> **Expensive**
- Provisioning for **less than peak** capacity ==> **Business impact**
- Monitoring and scaling **manually** ==> **Risk/requires expertise**

## Aurora Serverless in a Nutshell

- Smooth scaling, no downtime.
- Scales as your load **increases/decreases**.
- RDS Data API allows integration with Lambda and other **AWS services**.
- Warm buffer pool is **replicated via shared distributed storage** so cutover to newly scaled instances is quick.
- Scaling is generally done if >70% CPU utilization or >90% connection utilization for a period of **five minutes**.
- A monitoring service actively monitors **all DB instances** in the back end for performance, resources, etc.
- You can choose to pause service or take capacity to zero to cut down costs when there's **no incoming traffic**.
- Excellent for intermittent, unpredictable **workloads**.

Back

Next

[Back to Main](#)



Linux Academy

# Back-End Data Persistence Layer: AWS Aurora Serverless

Course Navigation

Aurora Serverless Concepts

Serverless Defined

Section 2

Front-End Serverless Layer: S3 and CloudFront

Section 3

Application Layer: AWS Lambda Functions

Section 4

Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

Overview

What Is Amazon Aurora Serverless?

Aurora Serverless Concepts

Integrating Aurora Serverless with Lambda Functions

Monitoring Serverless Application Metrics

Section 6

Back to Main

Aurora Serverless: Behind the Scenes



User/App



Fleet of Managed Proxy Servers



Load Balanced Connection



DB Instance

Scaling



Warm Pool of DB Instances



Shared Distributed Storage b/w DB Instances

Aurora Serverless Service



Linux Academy

# Back-End Data Persistence Layer: AWS Aurora Serverless

Course Navigation

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

### Overview

### What Is Amazon Aurora Serverless?

### Aurora Serverless Concepts

### Integrating Aurora Serverless with Lambda Functions

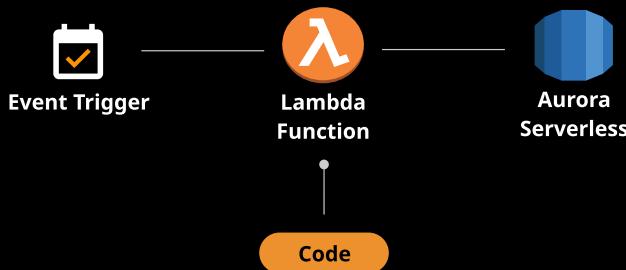
## Monitoring Serverless Application Metrics

Section 6

## Integrating Aurora Serverless with Lambda Functions

In almost all cases, you would want your database to be **restricted** and **only** accessible to your business logic applications or functions such as AWS Lambda.

### Integration Workflow



In our case, we will be using **Boto3**, which is an AWS Python SDK, in our Lambda code.

Points to note:

- RDS credentials (user/pass) are stored in Secrets Manager as a secret, so Lambda will require access to Secrets Manager to get the value.
- RDS will need access to RDSDatabaseService via the **rds-data:ExecuteStatement** action.

[Back](#)

[Next](#)

[Back to Main](#)



Linux Academy

# Back-End Data Persistence Layer: AWS Aurora Serverless

Course Navigation

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

### Overview

#### What Is Amazon Aurora Serverless?

#### Aurora Serverless Concepts

#### Integrating Aurora Serverless with Lambda Functions

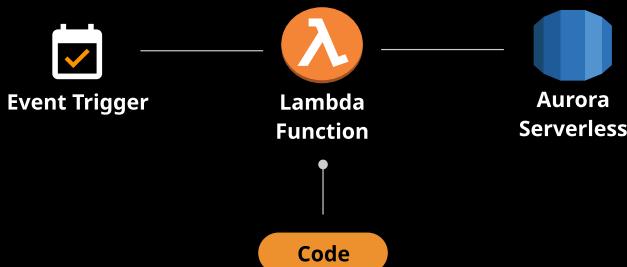
## Monitoring Serverless Application Metrics

Section 6

## Integrating Aurora Serverless with Lambda Functions

In almost all cases, you would want your database to be **restricted** and **only** accessible to your business logic applications or functions such as AWS Lambda.

### Integration Workflow



```
#NOT ACTUAL CODE
import boto3
client = boto3.client('rds-data')
.....
.....
def lambda_handler(event,context):
    response = client.execute_statement(database,
                                         resourceArn,
                                         secretArn,
                                         sql)
    return response
```

[Back to Main](#)



Linux Academy

# Monitoring Serverless Application Metrics

Course Navigation

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

## Monitoring Serverless Application Metrics

Section 6

## Metrics and CloudWatch Dashboard for Serverless

## Conclusion

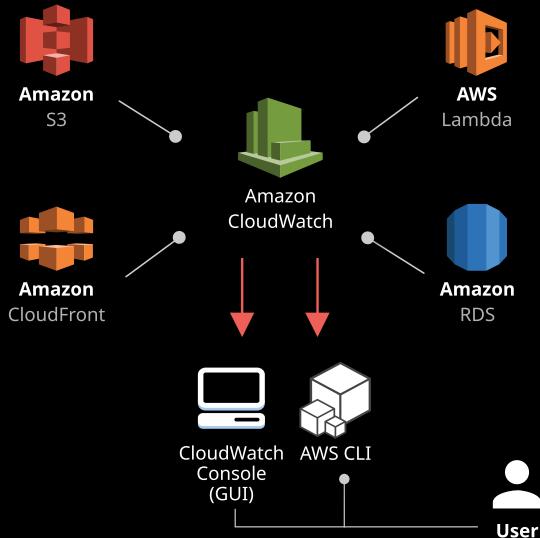
Section 7

Metrics and CloudWatch Dashboard for Serverless

## AWS CloudWatch and Metrics

### Serverless Metrics

- **Monitoring service** for AWS cloud
- CloudWatch provides **useful metrics** for all serverless services, such as S3, CloudFront, Lambda, and Aurora RDS
- Allows creating **customized dashboards**
- Integrates with other AWS services for **alarms and notifications**



Back

Back to Main



Linux Academy

# Monitoring Serverless Application Metrics

Course Navigation

What's Next?

## Serverless Defined

Section 2

## Front-End Serverless Layer: S3 and CloudFront

Section 3

## Application Layer: AWS Lambda Functions

Section 4

## Back-End Data Persistence Layer: AWS Aurora Serverless

Section 5

## Monitoring Serverless Application Metrics

Section 6

## Conclusion

Section 7

What's Next?

Get Recognized!

Back to Main

If this course piqued your interest in IAM , Lambda or even Python Boto SDK for AWS be sure to check out our courses:

- **AWS IAM Deep Dive**  
(<https://linuxacademy.com/cp/modules/view/id/180>)
- **Lambda Deep Dive**  
(<https://linuxacademy.com/cp/modules/view/id/204>)
- **Automating AWS with Lambda, Python and Boto3**  
(<https://linuxacademy.com/cp/modules/view/id/313>)

Connect with us and your fellow learners!

**Linux Academy community Slack**

<https://linuxacademy-community-slack.herokuapp.com>

**Linux Academy Community**

<https://linuxacademy.com/cp/socialize>



Linux Academy